

A Third Executor: Tapping the Integrated GPU via OpenCL

Matias Saavedra

Monday 15th June, 2026

Binary: branch feat/igpu-opencl (intended tag binary-v6-igpu) built on top of binary-v5-affinity (fc33e29). New translation unit oclkernel.cpp/.h; the CUDA and CPU paths are unchanged. Baseline is the same binary in gpuMode=1 (discrete GPU only).

Resumen

The renderer used two executor classes — the discrete GPU (CUDA) and a pool of CPU worker threads. This report adds a **third**: an OpenCL backend (oclkernel.cpp) that runs the FP64 Mandelbrot kernel on the *integrated* GPU (Intel UHD 630, on the i7-9750H die). `compute()` now dispatches by an `ExecKind` enum, both GPU backends share a result-copy helper, and the iGPU runs on its own `CalcThr` with the same largest-first affinity as the dGPU. Output is **byte-identical to the CUDA path** (0 pixel diff). The headline A/B (100 frames, 1920×1080 , `diffT` = 0.1, `save` = 1, 12 threads, 3 reps, on AC) measured a **wall win, not the predicted wash: dGPU+iGPU+10CPU 44.32 s vs. dGPU+11CPU 57.72 s**, -23.2% (disjoint ranges; -72.6% vs. the 161.6 s CPU floor) — and it wins *while surrendering a CPU worker*. The mechanism is not the predicted “GPU saturated”: the binding lane is the CPU pool (busiest thread $\approx$ wall in every config), and the second GPU strips total CPU compute from 622 s to 435 s (-30%), faster than removing a CPU thread raises per-thread load. The two GPUs self-balance (34.4 vs. 30.4 s kernel).

1. Setup and the Change Measured

Workload and machine.

As in reports 12 and 13: the same `spec.in` (100 frames, 1920×1080 , zoom to `maxIterations` = 10,000), i7-9750H (6c/12t) + GTX 1660 Ti Max-Q (*discrete* GPU, CUDA) + Intel UHD 630 (*integrated* GPU, OpenCL), 12 threads, `diffThreshold` = 0.1, `pixelSizeThresh` = 32,768, `save` = 1, on AC power. The decomposition is unchanged from `binary-v5-affinity`: **5,608 leaves** in every configuration — only *which* processor computes each leaf differs.

Three executors, one dispatch point.

The worker model was binary (GPU thread vs. CPU thread). It becomes a three-way `ExecKind`; the work queue still distinguishes only GPU-class (largest-first) from CPU (smallest-first), so `WorkQueue` is untouched. The branch lives entirely in `compute()`’s dispatch and the new translation unit:

Código 1: mandelregion.cpp:187–193 — `compute()` picks a front-end by executor kind.

```

1 if (kind == EXEC_CUDA)
2   hostFE (upperX, upperY, lowerX, lowerY, pixelsX, pixelsY, &h_res, &pitch,
3         ↪ MAXGRAY);
4 else // EXEC_IGPU
5   oclFE (upperX, upperY, lowerX, lowerY, pixelsX, pixelsY, &h_res, &pitch,
6         ↪ MAXGRAY);
7
8 commitGPUResult (h_res, pitch / sizeof (int),
9                 kind == EXEC_CUDA ? "GPU" : "iGPU");

```

oclFE mirrors the CUDA hostFE signature exactly, so the only difference between the two GPU paths is the function called. Both then run the identical copy-and-metrics loop, factored out of the old CUDA branch into `commitGPUResult` (so the discrete path is byte-for-byte the same operations it was, only relabelled):

Código 2: `mandelregion.cpp:137–160` — the shared GPU result-copy, used by both backends.

```

1 void MandelRegion::commitGPUResult (unsigned int *h_res, int pitchInts,
2                                     const char *backend) {
3   // ... for each pixel: read h_res[j*pitchInts + i], accumulate iter/inset
4   //   metrics, write colormap[color] (or black for in-set) into the QImage.
5   //   Identical to the previous inline CUDA copy; only the label differs.
6 }

```

The OpenCL backend (`oclkernel.cpp`).

Compiled by `g++` (not `nvcc`) against CUDA's bundled CL headers and linked through the ICD loader (`-lOpenCL`). It exposes three calls that parallel `kernel.cu`: `OCLmemSetup` / `oclFE` / `OCLmemCleanup`. The kernel is a line-for-line port of `mandelKernel`, embedded as a string so there is no `.cl` file to locate at runtime:

Código 3: `oclkernel.cpp:62–85` — the FP64 Mandelbrot kernel, embedded.

```

1 #pragma OPENCL EXTENSION cl_khr_fp64 : enable
2 int diverge(double cx, double cy, int MAXITER) {
3   int iter = 0; double vx = cx, vy = cy, tx, ty;
4   while (iter < MAXITER && (vx*vx + vy*vy) < 4.0) {
5     tx = vx*vx - vy*vy + cx; ty = 2.0*vx*vy + cy; vx = tx; vy = ty; iter++;
6   }
7   return iter;
8 }
9
10 __kernel void mandelKernel(__global unsigned int *out, double upperX, double
11 ↪ upperY,
12 ↪ double stepX, double stepY, int resX, int resY, int MAXITER) {
13   int myX = get_global_id(0), myY = get_global_id(1);
14   if (myX >= resX || myY >= resY) return;
15   out[myY*resX + myX] = (unsigned)diverge(upperX + myX*stepX, upperY - myY*
16 ↪ stepY, MAXITER);
17 }

```

oclFE sets the kernel arguments, launches a 2D NDRange rounded up to whole

16 × 16 work-groups (the in-kernel bounds check drops the overhang, exactly as the CUDA grid/block guard does), reads the region back, and times the kernel and the read with OpenCL profiling events — the direct analogue of the CUDA-event timing in `hostFE`:

Código 4: `oclkernel.cpp:264–271` — launch + readback, device-timed by events.

```
1 clEnqueueNDRangeKernel(queue, kernel, 2, NULL, global, local, 0, NULL, &evKernel
   ↪ );
2 clEnqueueReadBuffer(queue, d_res, CL_TRUE, 0, resX*resY*4, h_res, 0, NULL, &
   ↪ evRead);
3 // CL_PROFILING_COMMAND_START/END on evKernel, evRead -> per-region
   ↪ kernel/read ms
```

Device selection and the per-thread context.

`OCLMemSetup` enumerates every OpenCL GPU device and picks the *integrated* one: it prefers an Intel-vendor GPU, accepts any non-NVIDIA GPU as a fallback (the discrete card is already driven through CUDA — running OpenCL on it too would be pointless contention), honours a `MANDEL_OCL_DEVICE` override, and aborts with install guidance if none is found. Because an OpenCL context is best used from one thread, the iGPU worker builds and tears down its own context, queue, program and buffers *inside its own* `run()` — file-scope state touched by exactly one thread, the same single-thread assumption `kernel.cu` makes:

Código 5: `main.cpp:172–189` — the iGPU worker owns its OpenCL state for its lifetime.

```
1 if (kind == EXEC_IGPU) OCLMemSetup (resX, resY);    // build context on this
   ↪ thread
2 bool isGPU = (kind != EXEC_CPU);                    // both GPU classes pull
   ↪ largest
3 while ((t = que->extract (isGPU)) != NULL) { t->examine (*que, kind); delete t; }
4 MandelRegion::printCPUSummary();
5 if (kind == EXEC_IGPU) OCLMemCleanup ();
```

The CLI.

`arg3` becomes a backend bitmask, `gpuMode`: bit 0 = discrete GPU (CUDA), bit 1 = integrated GPU (OpenCL). So 0 = CPU only, 1 = dGPU (the historical default, preserved), 2 = iGPU, 3 = both. In dual mode thread 0 is the dGPU (main thread, wrapped by `CUDAMemSetup/Cleanup`) and the iGPU runs on worker thread 1.

The runtime gotcha.

The UHD 630 (Coffee Lake, *Gen9.5*, 8086:3e9b) is *not* served by the current `intel-compute-runtime` (NEO ≥24.x dropped Gen8–Gen11; it loads but prints `FATAL: Unknown device`). The Gen9.5 part needs the *legacy* runtime (`intel-compute-runtime-legacy1` on Arch). With it installed, `clinfo` reports the device with `cl_khr_fp64` present, so the double-precision kernel builds.

Measurement.

A four-config wall A/B at 12 threads, 3 reps each, all on one binary in one thermal state so the comparison is disjoint (the discrete-GPU baseline is re-measured here, not reused from reports 05/12).¹ Raw data: `experiments/20-igpu-openc1/`.

2. Results

2.1. Wall Time — The Win

Tabla 1: Wall time by backend mode (12 threads, `diffT` = 0.1, `save` = 1, 3 reps). Δ is relative to `dGPU+11CPU`, the prior best. The mode-1 and mode-3 ranges are *disjoint* (every mode-3 rep beats every mode-1 rep), so -23.2% is signal, not noise.

Config	Mode	Wall mean (s)	Range (s)	Δ vs. mode 1
CPU12	0	161.64	159.87–163.03	+103.92 (+180%)
iGPU+11CPU	2	81.72	81.06–82.75	+24.00 (+41.6%)
dGPU+11CPU	1	57.72	57.34–58.01	— (baseline)
dGPU+iGPU+10CPU	3	44.32	44.14–44.62	−13.40 (−23.2%)

¹ The `save=1` PNG writes raise the absolute wall ~ 5 s above the `save=0` numbers of reports 08–13, so the 57.72 s baseline here is *not* comparable to report 12’s 49.86 s. A separate `save=0` cross-check (`experiments/20-igpu-openc1/ab_save0.csv`; alternating mode 1/mode 3, 3 reps) confirms the gap is `save` + thermal, not a regression: it reproduces report 12’s `save=0` baseline at mode 1 (**49.24 s** vs. 49.86 s) and gives mode 3 at **38.19 s** (-22.4% , a new best, -23.4% below report 12). The win is identical in both regimes; the CUDA path is byte-identical to `binary-v5-affinity`.




Figure 1: End-to-end wall by backend mode (bars, left axis; red line = % of the dGPU+11CPU baseline, right axis). The green bar is the new configuration: adding the iGPU drops the wall to 76.8% of the current best. Error bars are min/max over 3 reps — invisibly small. Note the iGPU *alone* (mode 2) already halves the CPU floor.

2.2. Where the 5,608 Leaves Ran

Tabla 2: Per-executor breakdown (rep 1; the reps agree to < 1%). Leaves sum to 5,608 in every row (decomposition-identical). “CPU/thread” is total CPU compute over the CPU-thread count; it tracks the wall to ~98%, identifying the CPU pool as the binding lane.

Config	dGPU (reg/s)	iGPU (reg/s)	CPU reg	CPU thr·s	CPU/thr (s)	out- liers
CPU12	0	0	5,608	1,898	158.2	4
iGPU+11CPU	0	2,960 / 65.8	2,648	898	81.6	0
dGPU+11CPU	3,447 / 42.9	0	2,161	622	56.6	0
dGPU+iGPU+10CPU	2,915 / 34.4	1,422 / 30.4	1,271	435	43.5	0



Figure 2: The 5,608 leaves split by executor class. Every bar reaches the same dashed total (the decomposition never changes); only the assignment shifts. In the winning mode 3 the CPU pool (orange) is squeezed to 1,271 leaves — the two GPUs between them absorb 4,337 — and the iGPU (green) claims 1,422, a third of the GPU work.

2.3. Correctness

The iGPU output is **byte-identical to the CUDA path**: rendering one frame on each and comparing gives **0** differing pixels (ImageMagick AE). Both GPUs differ from the CPU render by the *same* 25,112 pixels (8.2%) — the pre-existing FP-rounding difference between the GPU compilers’ fused multiply-add and the host `g++` build, flipping the iteration count by ± 1 in the high-gradient colour bands where the `i%256` colormap makes a ± 1 flip a large colour delta. It is not an iGPU artifact; the two GPUs agree exactly.

3. Analysis

3.1. The win is real, large, and disjoint

The mode-1 and mode-3 ranges do not overlap (Table 1: 57.34–58.01 s vs. 44.14–44.62 s, a ~ 13 s gap), and the within-config spread is under 0.5 s. The -23.2% is therefore unambiguous signal. It is the largest single-step wall improvement in the series — GPU affinity (`12-gpu-affinity.tex`) moved the wall -3.6% ; this moves it -23.2% — and it is achieved while the configuration *gives up a CPU worker* (10 vs. 11), so the integrated GPU is worth far more than the eleventh CPU thread it displaces.

3.2. The mechanism: the CPU pool is the binding lane, and the iGPU shrinks its work

The wall is the busiest CPU thread.

In every accelerated config the CPU pool is tightly balanced and each CPU thread is busy for essentially the whole run: total CPU compute over the CPU-thread count equals the wall to $\sim 98\%$ (Table 2, “CPU/thr” vs. Table 1: 56.6 vs. 57.7 s; 43.5 vs. 44.3 s; 81.6 vs. 81.7 s). So to first order wall $\approx$ (total CPU compute)/(#CPU threads).

Adding the iGPU lowers the numerator faster than the denominator.

Mode 3 removes one CPU thread (11 $\rightarrow$ 10, which alone would *raise* per-thread load by 10%) but the two GPUs together strip total CPU compute from **622 s to 435 s** (-30% , -187 s) by absorbing 1,166 more leaves off the pool (CPU leaves 2,161 $\rightarrow$ 1,271, Fig. 2). Net per-thread CPU load falls 56.6 $\rightarrow$ 43.5 s (-23%), and the wall follows. This is why the prediction failed (next subsection): the lever is CPU-work *removal*, and a second GPU removes more of it than a single GPU could.

Why one GPU could not do this alone.

The dGPU in mode 1 is busy only 42.9 s of kernel time within a 57.7 s wall — it has headroom but stops pulling useful work once the *large* regions are exhausted, because below a size threshold the GPU’s per-region launch + synchronous-copy overhead (reports 01, 15) makes small regions cheaper on a CPU thread. A second hungry GPU lane raises the size at which “CPU is better” bites, so the GPUs collectively claim more of the medium-large regions that would otherwise pin CPU threads.

3.3. The prediction was wrong — and why

The pre-measurement call (recorded on the branch) was “marginal-to-negative: FP64-weak iGPU thermally coupled to the CPU pool, which is the binding wall.” The direction was right about *what* binds (the CPU pool) but wrong about the *consequence*. The error was assuming a third executor can only *rebalance* a binding CPU pool (which cannot help a balanced pool — the lesson of **11-priority-queue.tex**). In fact the iGPU does not rebalance the CPU pool; it *removes work from it*, which is exactly the work-reduction lever that reports 09 and 12 named as the only way past the post-affinity floor. The feared thermal/bandwidth contention is real but second-order here: the dGPU is discrete silicon (uncoupled), and the iGPU’s contention with the CPU pool is dwarfed by the 187 s of CPU compute it offloads.

3.4. The two GPUs self-balance; the iGPU alone already halves the floor

Under largest-first affinity the two GPUs end mode 3 with **34.4 s (dGPU)** and **30.4 s (iGPU)** of kernel time — within 12% despite the iGPU’s $\sim 1.8\times$ higher per-region cost (21.4 vs. 11.8 ms/region) — because the faster device simply pulls more regions. Neither GPU is the bottleneck (both sit well under the 44 s wall). Separately, the iGPU *as the sole accelerator* (mode 2) cuts the CPU-only floor nearly in half (161.6 $\rightarrow$ 81.7 s, -49%) by absorbing the interior outliers: CPU outliers (>5 s regions) drop from 4 to 0 in every accelerated mode (Table 2), including the 13.5 s frame-73 minibrot interior that pins a CPU thread in mode 0.

4. Recommendations

1. **Adopt `gpuMode = 3` (both GPUs) as the default on this machine.** It is a clean -23.2% ($57.72 \rightarrow 44.32$ s) with disjoint A/B ranges, decomposition-identical, byte-identical output, and contained to one new translation unit. Merge `feat/igpu-openc1` and tag `binary-v6-igpu`.
2. **Try strict dGPU-priority-on-the-largest.** Both GPUs currently pull from the same end, so the weaker iGPU can grab the single biggest outlier the dGPU would clear in $\sim 1/1.8$ the time. Reserving the very largest region for the dGPU (iGPU takes second-largest) may shave the tail further; expected to be small (the GPUs already self-balance to 12%) but it is a pure scheduling change in `WorkQueue`.
3. **Run the full Fig. 11.15 thread-scaling sweep in mode 3.** This report fixed 12 threads; the dGPU-vs-dGPU+iGPU scaling curves (`GPU+kCPU`) would show where the iGPU’s contribution saturates as CPU workers are added, comparable to `experiments/05-fig1115-postfix`.
4. **Generalise off this laptop with care.** The win is conditional on the CPU pool being the binding lane at `diffT=0.1` with `save=1`; on a machine where the dGPU saturates first, a second weak GPU would help less. The mechanism (CPU-work removal), not the magnitude, is what transfers.

5. Conclusion

Adding an OpenCL backend for the integrated GPU — a third `ExecKind` that shares the dGPU’s largest-first affinity and the CPU’s work queue, dispatched at one point in `compute()` — cuts the end-to-end wall by 23.2% ($57.72 \rightarrow 44.32$ s, disjoint ranges), and by 72.6% against the CPU-only floor. The configuration wins despite trading a CPU worker for the iGPU thread. The mechanism is measured, not assumed: the CPU pool is the binding lane (busiest thread $\approx$ wall), and the second GPU strips total CPU compute by 30% ($622 \rightarrow 435$ s), more than offsetting the lost worker. This refutes the pre-measurement “marginal-to-negative” prediction, which mistook a *work-removal* lever for a rebalancing one. The two GPUs self-balance to within 12%, the iGPU output is byte-identical to CUDA, and the iGPU alone halves the CPU floor. Scheduling was exhausted at -3.6% (`12-gpu-affinity.tex`); adding hardware that removes CPU work is what moved the wall again.